

# Linux Basics for Hackers - Notes

## 1) BINARIES

Located in -> /usr/bin and/or /usr/sbin

## 2) LINUX FILE SYSTEM

'/' -> The actual root system. The top most.

Inside it ->

/boot - Kernel image

/home - user dir

/proc - view of internal kernel data

/dev - special device files

/sbin - binaries

/root - SuperUser's Home Dir (different from '/')

/etc - Sys config

/mnt - GenPurpose Mount point

/sys - Kernel's view of HW

/bin - also binaries

/lib - libraries

/usr -> /sbin, /bin, /lib (more of the same stuff) /media - for ejectable media

## 3) cd Command

use '..' to move up 1 level

'.. ..' for 2 levels & '.. .. .' for 3 levels and so on

## 4) Search-based commands

\$ locate 'find\_this' -> finds all occurrences

\$ whereis 'module\_name' -> finds all BINARIES of the target (usually with man pages)

\$ which aircrack-ng -> finds the binary file located in the PATH variable of the system

\$ find -directory -option -targetExp -> finds literally everything [Eg: find /etc -type f -name apache2] -- altho apache2.\* will find all file extensions but first name as apache2

Lastly, grep to filter

\$ ps aux | grep apache2 -> will filter from all auxilliary processes containing apache2

## 5) 'cat' is versatile

\$ cat file\_name -> will spill the file contents.

\$ cat > file\_name -> will let you write in it BUT WILL REPLACE ALL EXISTING DATA.

\$ cat >> file\_name -> will actually let you append the text you enter.

## 6) Renaming doesn't exist in Linux

So we use

\$ mv newfile newfile2 -> to essentially rename the file

## 7) REMOVING A DIRECTORY

Only the directory : \$ rmdir

When that fails : \$ rm -r (recursively delete everything in it)

## 8) TEXT MANIPULATION

head and tail ->

\$ head file\_name -> first 10 lines

\$ head -n file\_name -> n number of lines from the start

\$ tail -> for bottom lines (+ specialized with a count)

\$ nl path\_to\_file\_or\_file\_name -> will number all the lines. // cat can and should be clubbed with grep as and when needed. Eg: cat snort.conf | grep output

**Sed command** - for find and replace ->

\$ sed s/search\_term/replace\_term/g path\_to\_filename > newfile\_name

-> 's/' will find the term, '/g' is for replacing globally. Rest is elementary.

-> Removing the '/g' will only replace the first occurrence.

Adding a number there can limit the number of occurrences to be changed. '/3' will only replace the first 3.

Eg: sed s/mysql/MySQL/g /etc/snort/snort.conf > snort2.conf

\$ more file\_name -> offers a scroll-able page if the file is too big.

\$ less file\_name -> "less is more" -- offers a filter to search for the term should you need to -- use the '/' key. Still scroll-able but with better functionality.

## 9) NETWORKING

'loopback' addr -- same as 'localhost' = 127.0.0.1

iwconfig - check for wireless adapter info -- good for getting power, the mode [monitor,

managed, promiscuous] etc.

Changing info -> MAC or IP Addr :

```
$ ifconfig eth0 192.168.181.155
```

-> modifies what your router sees to redirect packets

### **Changing Netmask +/- Broadcast :**

```
$ ifconfig eth0 192.168.181.155 netmask 255.255.0.0 broadcast 192.168.1.255
```

-> Netmask is the subnet mask -- determines the portion of the IP Addr to the NW and which refers to the host. Here, first 2 octets (16 bits) represent the network and last 2 show the hosts within that network.

-> Broadcast is the addr used to send packets to all hosts on the same network segment. Default (bcz of subnet) would become 192.168.255.255 but here overridden to 192.169.1.155 -- Basically packets sent to this will be broad-casted on that specific sub-network.

### **MAC Spoofing :**

Take down the interface, change the Addr, restart

```
$ ifconfig eth0 down > ifconfig eth0 hw ether 00:11:22:33:44:55 > ifconfig eth0 up
```

### **Assigning new IP via DHCP Server :**

Server runs of '*dhcpcd*' - the daemon. Requested via '*dhclient*'. Requires a DHCP assigned IP addr. (Note : '*dhclient*' is for Debian, different for all other distros.)

```
$ dhclient eth0
```

-> DHCPDISCOVER req is sent by this command and then receives an offer from the DHCP i.e DHCPOFFER. Now, ifconfig will show a diff IP addr as given by the DHCP Server.

### **Manipulating DNS :**

Use 'dig' :

Directly pass-on the domain and add the 'ns' tag to make the domain as the nameserver itself. 'mx' will fetch the mail-exchange server.

```
$ dig hackerarise.com ns OR $ dig hackerarise.com mx
```

-- Some Linux servers use BIND (Berkeley Internet Name Domain) which is just a fancy name for DNS.

### **Changing your DNS Server :**

Edit a file stored in '/etc/resolv.conf'. There you will see the domain, search & nameserver fields. Swap the values here to switch your DNS server.

Other method to do the same (although this cleanly overwrites the file's content) ->

```
$ echo "nameserver 8.8.8.8" > /etc/resolv.conf
```

## Mapping your own IP Addr :

'hosts' file located in /etc/hosts. Useful for hijacking a TCP connection on your LAN to direct traffic to a malicious webserver by using a tool like *'dnsspoof'*

Usually this file has a mapping for your localhost only BUT you can map any website to any IP Address. Eg : "192.168.181.131 bankofamerica.com". Decent for local network attacks.

[dnsspoof & Ettercap can be used]

## 10) HANDLING SW PACKAGES

### Search for package in local repo:

\$ apt-cache search *'keyword'*

-> Eg: \$ apt-cache search snort

### Install, Remove, Purge, Update, Upgrade :

\$ apt-get install --name--

\$ apt-get remove --name--

\$ sudo apt-get update

\$ sudo apt-get upgrade

\$ apt-get purge --name--

(purge removes the config as well)

or use a package manager like *'synaptic'* or *'gdebi'* like a normie.

### Adding Repos to Sources.list file :

\$ mousepad /etc/apt/sources.list -> will open the list

Categories : main (OSS), universe (community maintained OSS), multiverse (SW restricted by copyright), restricted (proprietary device drivers), backports (packages from later releases)

Format : "deb http:// ----- --package\_name-- main non-free contrib" etc.

## 11) FILE DIRECTORIES & PERMISSION

r,w,x -> read, write and execute

### Granting ownership :

\$ chown *username file-path-name* -> Provides the ownership of that file to that user

\$ chgrp *group-name package-or-module-name* -> Provides a user-group access to that module

### Checking Permissions :

\$ ls -l *file-or-path-to-it* -> will lay down the whole sheet. The type, permission on the file for

owner/groups/users, number of links, the owner, size in bytes, creation/mod date & its name.

Eg: "drwxr-xr-x" vs "-rw-r--r--". First letter denotes directory if 'd' or file if empty dash. Followed by the permission values for 3 groups i.e owner then group then other\_users. Hence we observe 3 values at a time. Dash means no permission ofc.

In '-rw-r--r--' -> File, owner has read/write, group and other users only have read permissions.

2nd way : There is a proper calculation done in Octal terms as well.

```
001 : 1 : --x
010 : 2 : -w-
011 : 3 : -wx
100 : 4 : r--
101 : 5 : r-x
110 : 6 : rw-
111 : 7 : rwx
```

Total RWX is 7. Since we have 3 sets of permissions, giving a full read+write+execute permission to everyone, for example, would look like ->

```
$ chmod 777 hashcat.hcstat
```

3rd way : UGO Syntax

Here, '-' removes a permission, '+' adds and '=' sets a permission.

Eg: Remove the write (w) permission from user on a file

```
$ chmod u-w hashcat.hcstat -> Now -rw-r--r-- becomes -r-xr-xr--
```

Or for user and other users at once

```
$ chmod u+x, o+x hashcat.hcstat
```

Now, you can set execute permission for yourself on a newly downloaded tool/script bcz by default Linux won't set it

```
$ chmod 766 some_new_tool -> grants us (the owner) all permission including execute -- and everyone else only R/W permissions.
```

### **Masking can be done :**

```
$ umask 007 -> set it so only the user and members of the user's group have permissions.
```

### **Special Permissions :**

SUID - set user ID & SGID - set group ID

1) Granting Temp Root w/ SUID

/etc/shadow contains all user's password -- requires root privileges to execute. SUID requires an additional bit before the permission bit. So 644 becomes 4644 i.e

```
$ chmod 4644 file_name
```

## 2) Granting the Root user's Group permissions SGID

SGID works differently. Someone without execute permission can execute a file if the owner belongs to the group that has the permission to execute that file. When the bit is set on a directory -- **the ownership of new files created in that directory goes to the directory's creator's group rather than the file creator's group.**

`$ chmod 2644 file_name [ SGID bit is represented by 2 and SUID uses 4]`

### Privilege Escalation :

One way - exploit the **SUID Bit** in the system. Eg: Scripts that need to change the password usually come with the SUID bit set already. Use that to gain temporary root priv - then do something shady like getting the file at /etc/shadow.

To proceed -> Use commands like 'find' to find the files and see their bit. Example :

`$ find / -user root -perm -4000`

Kali now starts at the top of the filesystem (because of '/') and looks everywhere below this -- the file that are owned by 'root' & specified with 'user root' + have the SUID bit set (-perm -4000). This command will give an output like ->

`/usr/bin/chsh ; /usr/bin/gpasswd; /usr/bin/pkexec; /usr/bin/sudo; /usr/bin/passwd,.. etc.`

Navigating to this directory, and observing, let's say "sudo" using ls-lalh, you will see ->

`-rwsr-xr-x root root 140944 date sudo`

Here, the 's' in place of 'x' determines the SUID bit. Logically, anyone who runs the *sudo* file has the priv of a root user -- which becomes an attack vector IF an application -- which needs access to /etc/shadow file to successfully complete their task -- can be hijacked.

## 12) PROCESS MANAGEMENT

To view - use -> `$ ps`

Every process ofc has a PID or process ID. You can use -> `$ kill PID_value` to kill any process.

Issue ? 'ps' command won't give you much info either ways. We have another command for that ->

`$ ps aux`

It shows the USER, PID, %CPU, VSZ, RSS, TTY, STAT, START, TIME & COMMAND.

### Filtering by Process Name

For instance, try running *msfconsole* command to have its process running. Then use *grep* to filter it. This way you can filter all the processes running/attached to it.

`$ ps aux | grep msfconsole`

You might see a few, such as the attached DB running, the ruby script, etc, and finally the program itself.

We also have commands like "*top*" to monitor the processes sorted by their resource usage. It's active i.e refreshes on its own (every 3-4 seconds)

## Managing Processes

We can alter the affinity/priority of any process by using the "*nice*" command (by passing a numeric value to its argument '-n'). Kernel always has the final say, we're just suggesting.

The value ranges from -20 to +19. Sadly, **HIGH value means LOW priority and vice versa**. So -20 is most likely to receive priority, 0 is default ofc, and +19 is least likely.

Usually, any process inherits the *nice* value of its parent process.

Unsurprisingly, you can alter the priority by using the "*renice*" command.

DIFFERENCE !!

**nice is relative.** Its adds/subtracts the priority value given what you pass to it. A process with a priority of 15, when asked 'nicely' to be -10 will have a priority of 5 now. OR when asked to be +5, it will now be 20. 'nice' can use the process via its location as well.

**renice is absolute.** Requires a fixed value b/w -20 and +19. BUT it sets the process to that level, cuz you've altered the deal and it prays you don't alter it any further. It also requires the PID.

Examples :

```
$ nice -n 10 /bin/some_slow_process [lowers it]
```

or

```
$ nice -n -9 /bin/some_slow_process [improves it]
```

and

```
$ renice 20 6996 [6996 is the PID of some_slow_process, and 20 is setting it]
```

NOTE: 'top' can also be used to alter these values.

## Killing Processes

'*kill*' command is your friend. Just pass the PID and pass the required kill signal. There are 64 of them. Default is SIGTERM (n=15) i.e termination. Ofc they are optional. Use them as a flag arg while using kill command.

```
$ kill -n PID
```

Example - \$ kill -9 6887

Signal Interrupts for kill ->

SIGHUP (1) : Hangup - stops and then restarts with the same PID

SIGIN (2) : Interrupt - weak kill signal not guaranteed to work but does work mostly.

SIGQUIT (3) : Quit/Core dump - terminates but saves the process info in memory + inside pwd.

SIGKILL (9) : absolute kill signal. Forces the process to stop by sending the process's resources to a special device -- /dev/null

Basically : to restart - use '-1' ; for zombie/malicious - use '-9'

## **Running Processes in the Background**

Everything runs from within the shell and the shell waits for the task/command to run/finish. It waits for this whole sequence -- hence busy & won't allow any new commands. To prevent this we can essentially detach the process from the shell. Use the '&' right after the task.

```
$ mousepad someDoc &
```

## **Moving to foreground**

Use the 'fg' command followed by the PID. Fetch the PID if needed.

```
$ fg 1273
```

## **Process Scheduling**

Either use 'at' or 'crond'.

'at' is useful for scheduling a job to run once at some point in the future -- execute 1 or many commands in the future, passed with time as argument.

Eg: \$ msfconsole at 7:20PM June 13

OR \$ msfconsole at now + 20 minutes

// If you just write 'at' followed by a time, then the "at>" console will open asking you to map the file/process path for it.

'crond' is best for scheduling tasks to occur everyday/week/month etc. [SEPARATE CHAPTER LATER]

## **13) MANAGING USER ENVIRONMENT VARIABLE**